

# PredictingStocks\_X

Hair Parra

May 10, 2020

## Historical Stocks Data Anlalysis: Forecasting Closing Prices

### Loading packages

```
library(tidyverse)
library(tidyquant)
library(gridExtra)
library(tibbletime)
library(forecast)
library(itsmr)
library(here)
library(bbmle)
library(tseries)
library(fpp2)
library(ggthemes)
library(readr)
library(xts)
library(reshape)
require(timeDate)
library(png)
knitr::opts_chunk$set(comment=NA, tidy=TRUE)
```

### Loading the data

```
stocks_3M <- read_csv("../data_raw/stocks_data_3M.csv")
```

Parsed with column specification:

```
cols(
  Date = col_date(format = ""),
  Open = col_double(),
  High = col_double(),
  Low = col_double(),
  Close = col_double(),
  `Adj Close` = col_double(),
  Volume = col_double()
)
```

```
head(stocks_3M, 10)
```

```
# A tibble: 10 x 7
  Date      Open  High  Low Close `Adj Close` Volume
<date>    <dbl> <dbl> <dbl> <dbl>      <dbl>    <dbl>
1 2020-03-04 40.7  41.5 39.8 41.4      41.0 30022100
2 2020-03-05 40.2  40.5 39.3 39.6      39.2 30255900
3 2020-03-06 38    40.0 37.8 39.7      39.3 48605600
4 2020-03-09 36.9  39.6 36.3 38.0      37.6 61535300
5 2020-03-10 39.2  40.2 37.9 40.1      39.7 50536500
6 2020-03-11 39.0  39.2 36.4 37.0      36.7 63594300
7 2020-03-12 34.5  35.8 33    33.2      32.9 51855300
8 2020-03-13 35.2  37.7 33.3 37.6      37.3 53859600
9 2020-03-16 33.2  37.0 32.4 33.7      33.4 44211300
10 2020-03-17 34.7  36.2 33.6 35.5      35.2 41572400
```

## Data Preprocessing

Next, extract the columns of interest and convert into time series objects

```
stocks_3M_data <- select(stocks_3M, Date, Close) # extract cols
dates <- as.POSIXct.Date(stocks_3M_data$Date) # extract dates in POSIXct format
stocks_3M_data.ts <- xts(stocks_3M_data$Close,
                        order.by = dates) # 7600
str(stocks_3M_data.ts) # inspect the data
```


```
An 'xts' object on 2020-03-03 19:00:00/2020-04-30 20:00:00 containing:
Data: num [1:42, 1] 41.4 39.6 39.7 38 40.1 ...
Indexed by objects of class: [POSIXct,POSIXt] TZ:
xts Attributes:
NULL
```

## Inspecting the data

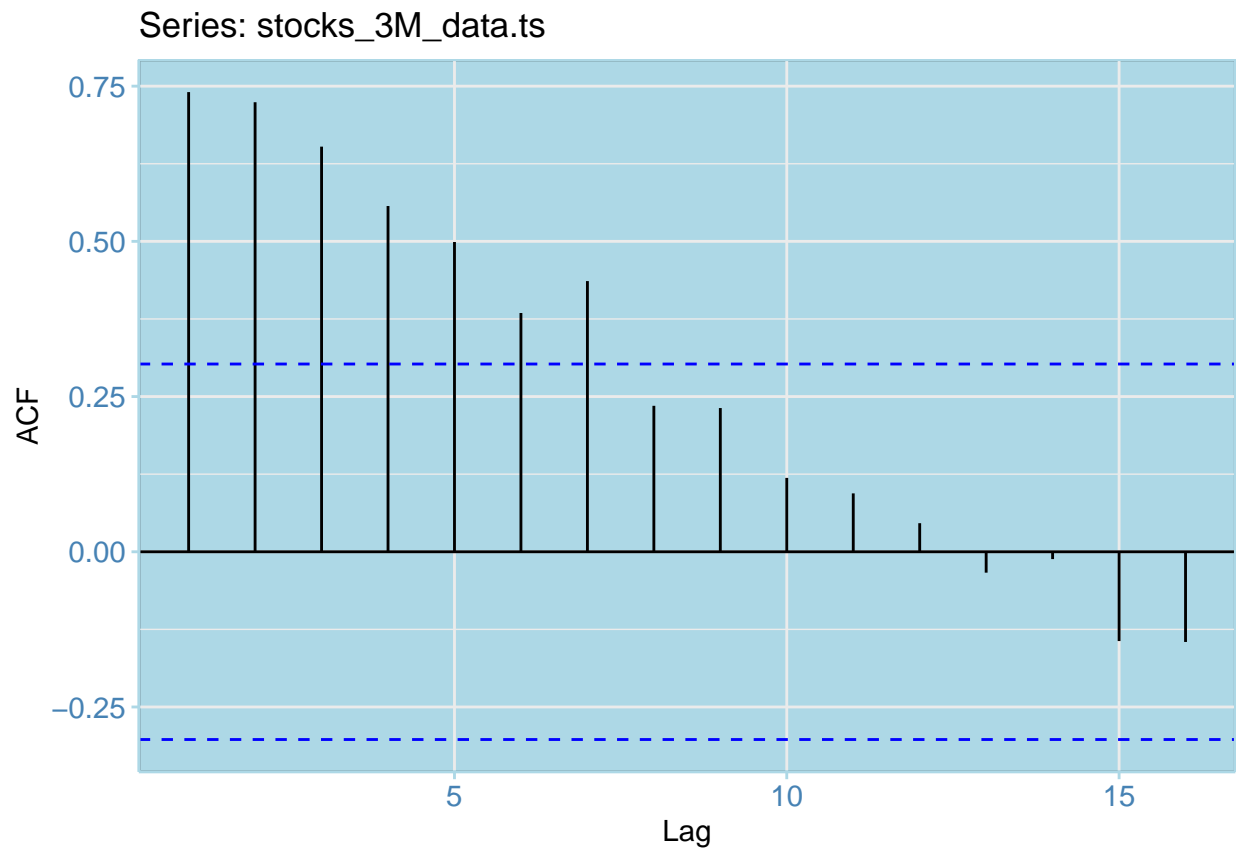
### Autoplot, ACF and PACF

```
# Plot the same white noise this time as lines
autoplot(stocks_3M_data.ts) +
  geom_line(colour="blue") +
  ggtitle("Stocks closing price historical data (3M)") +
  theme_stonks() + xlab("Date") + ylab("USD") + geom_point(color="black")
```

Stocks closing price historical data (3M)




```
# ACF  
ggAcf(stocks_3M_data.ts) + theme_stonks()
```



```
# PACF  
ggPacf(stocks_3M_data.ts) + theme_stonks()
```

Series: stocks\_3M\_data.ts



## Estimating the trend

```
# Estimate various trends
stocks_3M_linear <- tslm(ts(stocks_3M_data.ts)~trend)
stocks_3M_p5 <- tslm(ts(stocks_3M_data.ts)~trend + I(trend^2) + I(trend^3) + I(trend^4) + I(trend^5) )
stocks_3M_ma5 <- ma(ts(stocks_3M_data.ts), order=5) # moving average
stocks_3M_trends <- data.frame(cbind(Data=stocks_3M_data.ts, # stack in a dataframe
  Linear_trend=fitted(stocks_3M_linear),
  Poly_trend=fitted(stocks_3M_p5),
  Moving_avg5 = stocks_3M_ma5
))

# transform to xts objects
stocks_3M_linear <- xts(fitted(stocks_3M_linear), order.by = dates)
stocks_3M_p5 <- xts(fitted(stocks_3M_p5), order.by = dates)


# Plot all the trends together
autoplot(stocks_3M_data.ts, colour="original") + theme_stonks() +
  geom_line(aes(y=stocks_3M_linear, color="linear"),size=1) +
  geom_line(aes(y=stocks_3M_p5, color = "O(5) poly"), size=1) +
  geom_line(aes(y=stocks_3M_ma5, color = "ma21"), size=1) +
  scale_color_manual(values = c('original'= 'blue',
    'linear' = 'darkblue',
    'O(5) poly' = 'red',
```

```

      'ma21'= 'yellow')) +
labs(color = 'Trend fit') + ylab("USD") +
ggtitle("Different trend fits for the stocks data")

```

Warning: Removed 4 row(s) containing missing values (geom\_path).




```

# Detrend and show the de-trended series
stocks_3M_p5_xts <- xts(stocks_3M_p5, order.by = dates) # cast to xts
detrend_stocks_3M <- stocks_3M_data.ts - stocks_3M_p5_xts # subtract from original

# Plot the residuals
autoplot(detrend_stocks_3M) + theme_stonks() +
  ggtitle("De-trended Data ( O(5) trend)") +
  geom_hline(yintercept = 0, colour="black") +
  geom_point() + ylab("USD - trend") + geom_line(color="blue")


```

De-trended Data (  $O(5)$  trend)



The residuals look zero-trended.

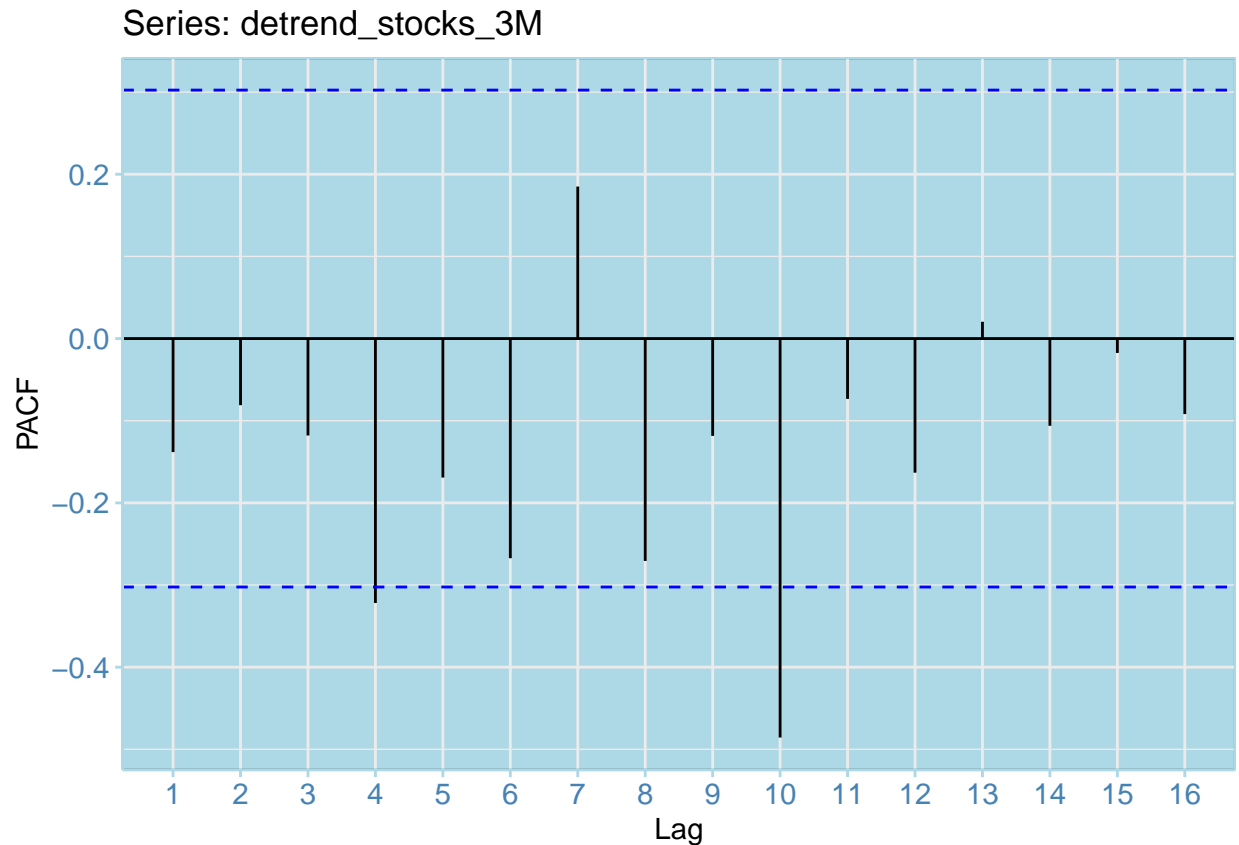
```
# ACF  
ggAcf(detrend_stocks_3M) + theme_stonks()
```



The ACF lags all , except for one fall within the 0.25 confidence bounds.

```
# PACF  
ggPacf(detrend_stocks_3M) + theme_stonks()
```





The PACF residuals mostly fall within the confidence bounds; however there seems to be some negative autocorrelation present across lags. However, from all the previous, there doesn't seem to be a strong seasonal component present.

## Train-test split & ARIMA fitting

We will now split the data into 32 training data points and 10 test data points. We will produce predictions and compare them to assess fit.

```
## train_test_split
detrend_stocks_3M_train <- stocks_3M_data.ts[1:(round(length(detrend_stocks_3M))-10)] # 32
detrend_stocks_3M_test <- stocks_3M_data.ts[(round(length(detrend_stocks_3M))-9):length(detrend_stocks_3M)]
str(detrend_stocks_3M_train)
```

```
An 'xts' object on 2020-03-03 19:00:00/2020-04-16 20:00:00 containing:
  Data: num [1:32, 1] 41.4 39.6 39.7 38 40.1 ...
  Indexed by objects of class: [POSIXct,POSIXt] TZ:
  xts Attributes:
  NULL
```

```
str(detrend_stocks_3M_test)
```

```
An 'xts' object on 2020-04-19 20:00:00/2020-04-30 20:00:00 containing:
  Data: num [1:10, 1] 42.5 40.5 41.8 41.6 42.5 ...
```

```
Indexed by objects of class: [POSIXct,POSIXt] TZ:
xts Attributes:
NULL
```

```
length(detrend_stocks_3M_train)
```

```
[1] 32
```

```
length(detrend_stocks_3M_test)
```

```
[1] 10
```

```
# Fit the ARIMA model on trian data
detrend_stocks_3M_arima_110 = auto.arima(detrend_stocks_3M_train,
                                         seasonal=TRUE,
                                         stepwise=FALSE,
                                         max.d = 2,
                                         ic = c("aicc", "aic", "bic") ,
                                         approximation=FALSE,
                                         trace=TRUE)
```

```
ARIMA(0,1,0) : 134.4715
ARIMA(0,1,0) with drift : 136.753
ARIMA(0,1,1) : 128.4086
ARIMA(0,1,1) with drift : 130.7508
ARIMA(0,1,2) : 129.5682
ARIMA(0,1,2) with drift : 132.1684
ARIMA(0,1,3) : 132.2178
ARIMA(0,1,3) with drift : 135.0296
ARIMA(0,1,4) : 134.9037
ARIMA(0,1,4) with drift : 137.877
ARIMA(0,1,5) : Inf
ARIMA(0,1,5) with drift : Inf
ARIMA(1,1,0) : 128.0794
ARIMA(1,1,0) with drift : 130.4957
ARIMA(1,1,1) : 129.8294
ARIMA(1,1,1) with drift : 132.4015
ARIMA(1,1,2) : 132.2178
ARIMA(1,1,2) with drift : 135.0353
ARIMA(1,1,3) : 134.9886
ARIMA(1,1,3) with drift : 138.0829
ARIMA(1,1,4) : Inf
ARIMA(1,1,4) with drift : Inf
ARIMA(2,1,0) : 129.5882
ARIMA(2,1,0) with drift : 132.1668
ARIMA(2,1,1) : Inf
ARIMA(2,1,1) with drift : Inf
ARIMA(2,1,2) : 134.9701
ARIMA(2,1,2) with drift : 138.0537
ARIMA(2,1,3) : Inf
ARIMA(2,1,3) with drift : Inf
```

```

ARIMA(3,1,0) : 132.1147
ARIMA(3,1,0) with drift : 134.9255
ARIMA(3,1,1) : 134.9128
ARIMA(3,1,1) with drift : 137.9894
ARIMA(3,1,2) : Inf
ARIMA(3,1,2) with drift : Inf
ARIMA(4,1,0) : 134.9736
ARIMA(4,1,0) with drift : 138.0206
ARIMA(4,1,1) : Inf
ARIMA(4,1,1) with drift : 141.3562
ARIMA(5,1,0) : 137.4053
ARIMA(5,1,0) with drift : 140.7549

```

Best model: ARIMA(1,1,0)

```
detrend_stocks_3M_arima_110
```

```
Series: detrend_stocks_3M_train
ARIMA(1,1,0)
```


```
Coefficients:
      ar1
      -0.4935
s.e.    0.1541
```

```
sigma^2 estimated as 3.237: log likelihood=-61.83
AIC=127.65 AICc=128.08 BIC=130.52
```

## Inspecting the residuals

```
checkresiduals(detrend_stocks_3M_arima_110)
```

Residuals from ARIMA(1,1,0)




Ljung-Box test

data: Residuals from ARIMA(1,1,0)  
 Q\* = 1.3259, df = 5, p-value = 0.9322

Model df: 1. Total lags used: 6

```
# Obtain dates and residuals
train_dates <- as.POSIXct.Date(stocks_3M_data$Date[1: (length(stocks_3M_data$Date) - 10 )])
arima_res <- xts(residuals(detrend_stocks_3M_arima_110),
                 order.by = train_dates)
```

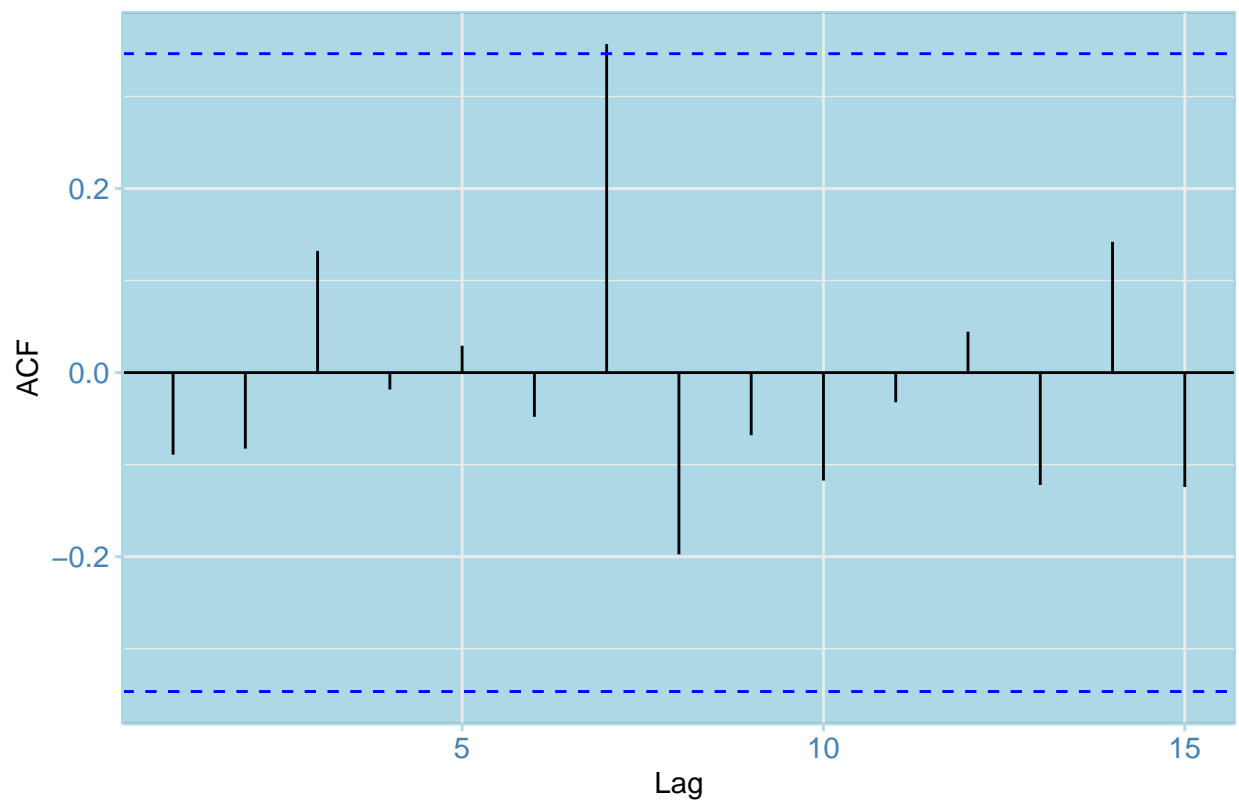
```
# Plot the residuals
autoplot(arima_res) + theme_stonks() +
  geom_point() + geom_line(color="blue") +
  geom_hline(yintercept = 0, colour="black")
```



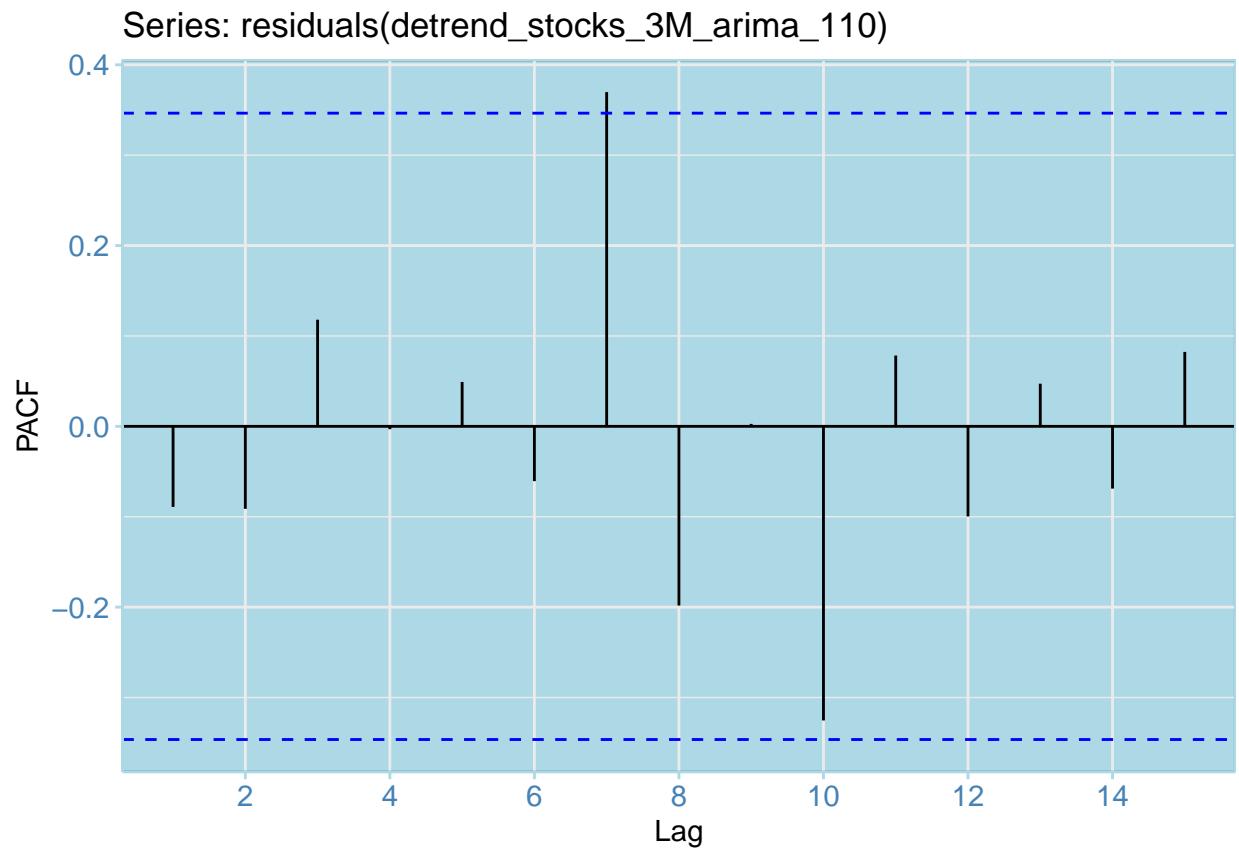
We see that perhaps around March 13, there could have been a possible outlier.

```
# ACF
ggAcf(residuals(detrend_stocks_3M_arima_110)) + theme_stonks()
```

Series: residuals(detrend\_stocks\_3M\_arima\_110)

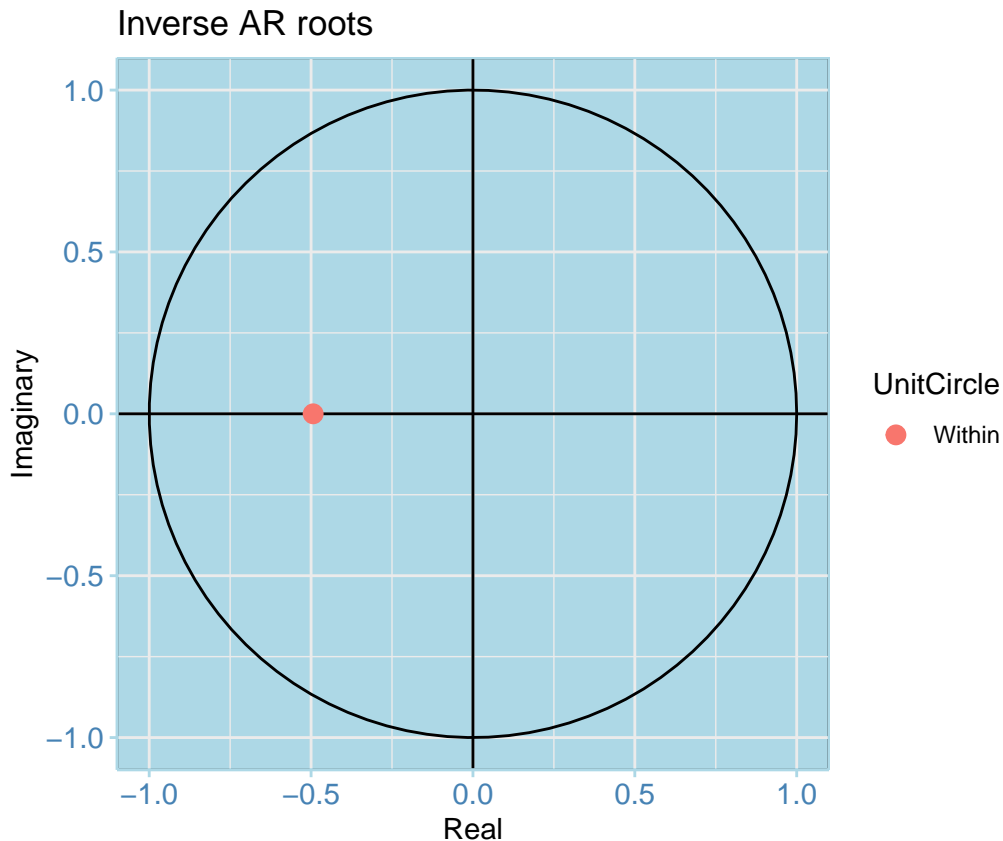


```
# PACF  
ggPacf(residuals(detrend_stocks_3M_arima_110)) + theme_stonks()
```



In both cases, the ACF and PACF points find within confidence bounds, with the exception of one. This one might be due to the possible outlier we had before.

```
# Inspect roots  
autoplot(detrend_stocks_3M_arima_110) + theme_stonks()
```



The roots of the AR(1) polynomial guarantee the process is stationary and causal, and of course, it is also invertible. We can also verify this by performing the ADF and KPSS tests for stationarity:

```
# Test with a bunch of different k's ? (bigger augmented versions)
detrrend_stocks_3M_arima_110_diff <- diff(residuals(detrrend_stocks_3M_arima_110), lag=1) # difference or
adf.test(detrrend_stocks_3M_arima_110_diff,k=1) # ADF
```

Warning in adf.test(detrrend\_stocks\_3M\_arima\_110\_diff, k = 1): p-value smaller than printed p-value

#### Augmented Dickey-Fuller Test

```
data: detrrend_stocks_3M_arima_110_diff
Dickey-Fuller = -7.2389, Lag order = 1, p-value = 0.01
alternative hypothesis: stationary
```

```
kpss.test(detrrend_stocks_3M_arima_110_diff) # KPSS
```

Warning in kpss.test(detrrend\_stocks\_3M\_arima\_110\_diff): p-value greater than printed p-value

#### KPSS Test for Level Stationarity



```
data: detrend_stocks_3M_arma_110_diff
KPSS Level = 0.055657, Truncation lag parameter = 2, p-value = 0.1
```

Reject -> stationary for the ADF, fail to reject -> stationary for the KPSS. Now we can proceed with the forecasting.

## Forecasting

### Obtaining model and trend forecasts

We will now forecast 10 observations from both the main model and the trend

```
detrend_stocks_3M_arma_110_forecasts <- forecast::forecast(detrend_stocks_3M_arma_110,h=10) # ARIMA(1
forecasted_trend <- forecast::forecast( stocks_3M_p5, h=10) # forecast 10 trend observations
```

### Model forecasts

Let's produce a table with the point forecast values along with the errors and confidence intervals for predictions

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2764801	41.99639	39.69064	44.30213	38.47005	45.52272
2851201	42.23504	39.65039	44.81969	38.28216	46.18793
2937601	42.11727	39.00740	45.22714	37.36114	46.87340
3024001	42.17539	38.74312	45.60766	36.92618	47.42459
3110401	42.14671	38.36444	45.92898	36.36223	47.93119
3196801	42.16086	38.08401	46.23771	35.92585	48.39587
3283201	42.15388	37.79058	46.51717	35.48079	48.82696
3369601	42.15732	37.53075	46.78389	35.08159	49.23305
3456001	42.15562	37.27741	47.03383	34.69504	49.61620
3542401	42.15646	37.04017	47.27275	34.33177	49.98115

```
# Show the table with errors
forecast_table = as.data.frame(forecast_table)
colnames(forecast_table) <- c("Point_Forecast","Lo80","Hi80","Lo95","Hi95","observed","errors")
forecast_table
```


	Point_Forecast	Lo80	Hi80	Lo95	Hi95	observed	errors
1	41.99639	39.69064	44.30213	38.47005	45.52272	42.54	-0.5436159
2	42.23504	39.65039	44.81969	38.28216	46.18793	40.55	1.6850426
3	42.11727	39.00740	45.22714	37.36114	46.87340	41.76	0.3572703
4	42.17539	38.74312	45.60766	36.92618	47.42459	41.62	0.5553886
5	42.14671	38.36444	45.92898	36.36223	47.93119	42.52	-0.3732934
6	42.16086	38.08401	46.23771	35.92585	48.39587	43.06	-0.8991408
7	42.15388	37.79058	46.51717	35.48079	48.82696	42.49	-0.3361264
8	42.15732	37.53075	46.78389	35.08159	49.23305	43.34	-1.1826776
9	42.15562	37.27741	47.03383	34.69504	49.61620	42.38	-0.2243795
10	42.15646	37.04017	47.27275	34.33177	49.98115	40.92	1.2364629

Extract the values as plain vectors for plotting: we paste this to a bunch of NA values to be able to plot all together.

```
predicts <- c(rep(NA,32),forecast_table$Point_Forecast)
predicts_Lo80 <- c(rep(NA,32),forecast_table$Lo80)
predicts_Hi80 <- c(rep(NA,32),forecast_table$Hi80)
predicts_Lo95 <- c(rep(NA,32),forecast_table$Lo95)
predicts_Hi95 <- c(rep(NA,32),forecast_table$Hi95)
```

## Producing the forecasts

```
# Plot the predictions + xlim(1.05e+08,1.09e+08) + ylim(32,45)
autoplot(detrend_stocks_3M_arima_110_forecasts) + theme_stonks()
```




Although this plot looks pretty, notice the scale is somehow very off! We will construct a better plot manually, by using the values we obtained before, so that we can see both the original data and the repdictions along with the confidence intervals like above.

```
autoplot(stocks_3M_data.ts, colour="orig") + theme_stonks() +
  geom_line(aes(y=predicts,colour = "predicted")) +
  geom_ribbon(aes(x=dates, ymin=predicts_Lo95,ymax=predicts_Hi95),fill="pink", alpha=.3) +
  geom_ribbon(aes(x=dates, ymin=predicts_Lo80,ymax=predicts_Hi80),fill="blue", alpha=.5) +
  scale_color_manual(values = c('predicted'= 'red','orig'='black')) +
```

```
ylab("USD") + xlab("Date") + geom_point() + geom_line(color="blue") +
geom_point() + ggtitle("Stocks predictions for 10 days")
```

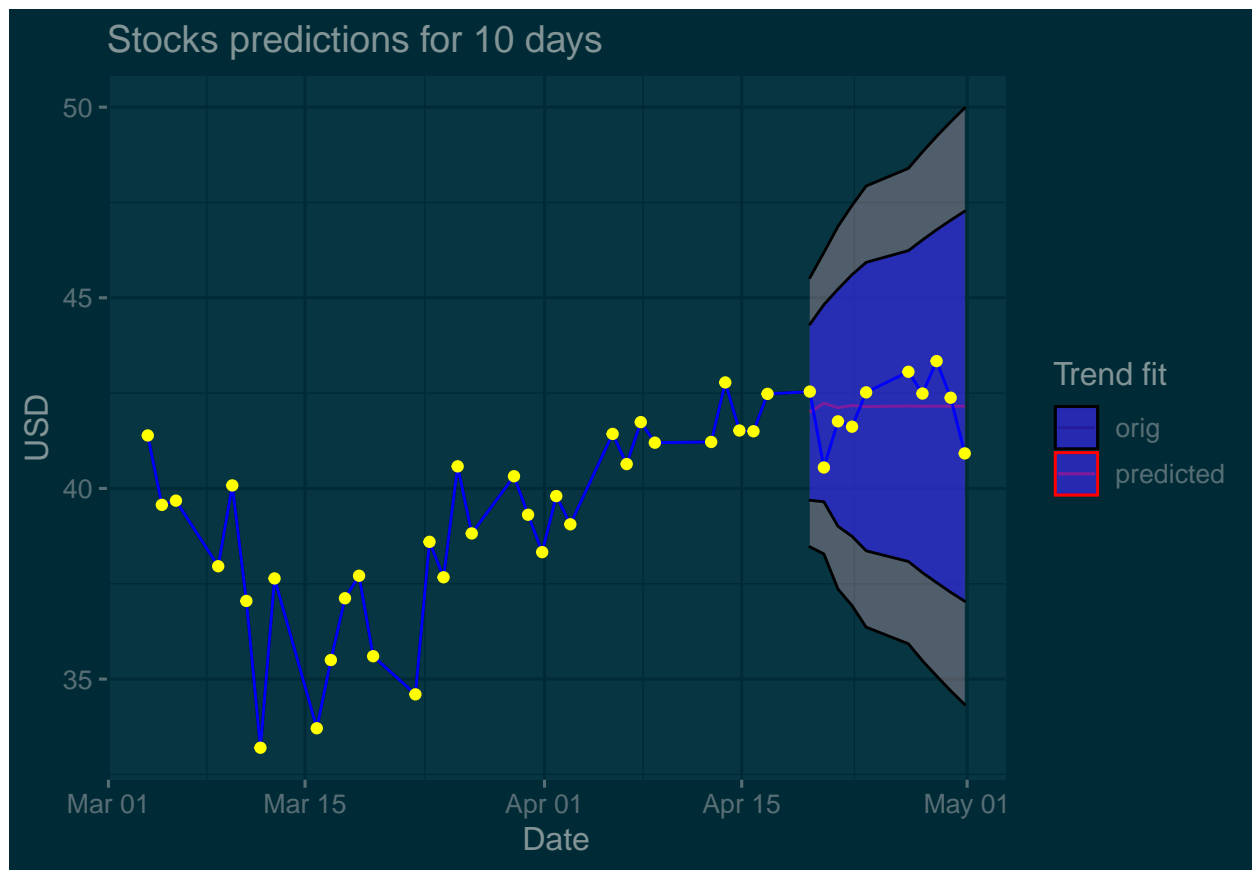
Warning: Removed 32 row(s) containing missing values (geom\_path).



```
autoplot(stocks_3M_data.ts, colour="orig") + theme_solarized_2(light = FALSE) +
scale_colour_solarized("blue") +
geom_line(aes(y=predicts,colour = "predicted")) +
geom_ribbon(aes(x=dates, ymin=predicts_Lo95,ymax=predicts_Hi95),fill="pink", alpha=.3) +
geom_ribbon(aes(x=dates, ymin=predicts_Lo80,ymax=predicts_Hi80),fill="blue", alpha=.5) +
scale_color_manual(values = c('predicted'= 'red','orig'='black')) +
labs(color = 'Trend fit')+ylab("USD") + xlab("Date") + geom_line(color="blue") +
geom_point(color="yellow") + ggtitle("Stocks predictions for 10 days")
```

Scale for 'colour' is already present. Adding another scale for 'colour', which will replace the existing scale.

Warning: Removed 32 row(s) containing missing values (geom\_path).



```
autoplot(stocks_3M_data.ts, colour="orig") + theme_hc(bgcolor = "darkunica") +
  scale_colour_hc("darkunica") +
  geom_line(aes(y=predicts,colour = "predicted")) +
  geom_ribbon(aes(x=dates, ymin=predicts_Lo95,ymax=predicts_Hi95),fill="pink", alpha=.3) +
  geom_ribbon(aes(x=dates, ymin=predicts_Lo80,ymax=predicts_Hi80),fill="blue", alpha=.5) +
  scale_color_manual(values = c('predicted'= 'red','orig'='black')) +
  labs(color = 'Trend fit')+ylab("USD") + xlab("Date") + geom_line(color="blue") +
  geom_point(color="yellow") + ggtitle("Stocks predictions for 10 days")
```

Warning in theme\_hc(bgcolor = "darkunica"): `bgcolor` is deprecated. Use `style` instead.

Scale for 'colour' is already present. Adding another scale for 'colour', which will replace the existing scale.

Warning: Removed 32 row(s) containing missing values (geom\_path).

